

MODERN DATA STRUCTURES

Columbia University
GR5072, Spring 2023
Wednesday, 6:00-8:00PM
332 Uris Hall

Instructor: Nick Anderson
Email: nla2121@tc.columbia.edu
Office Hours: Held by appointment
TA: Roxy Zhang
Email: rz2570@cumc.columbia.edu
Office Hours: TBD

I. Overview

This course is intended to provide some foundational day-to-day tools for a data scientist in industry. Participants will learn how to access and clean data, organize and test code using functions and classes, how to set up and publish their own python package, and use Github to keep code version-controlled and in sync with others. In addition, the course will enable participants to work with APIs, deal with data in a variety of formats, and access databases via SQL. Lastly, participants will receive an introduction to Amazon Web Services on how to store data, create databases, and parallelize computing.

The goal is to prepare students for their role as a data scientist (and incidentally a flavor of what would be expected of them in a typical data science interview.) Each week will have simple, moderate, and complex examples in class, with code to follow. Students will then practice additional exercises at home. No analysis or visualization (beyond just basic tables and plots to make sure everything was correctly organized) will be taught; and this will free up substantial time for the “nitty-gritty” of the data science workflow.

Some familiarity with Python, in particular with regards to the base functions (i.e., lists, tuples, dictionaries etc.) is assumed. If you have no prior Python experience, you can still take the class but you may have to spend some additional time studying Python as the class starts (The first week assigns a full review of base Python).

Course Website: Lecture materials, exercises, and HWs will be made available on GitHub only (except for the first and second weeks, as you may not yet have an github account / know how to use it). The GitHub repository is private. All enrolled students need to add their name to this [form](#) to be added.

II. Course Resources

Textbooks: There are no required textbooks for this course, but you will find these to be very useful in addition to the lectures and course readings:

- McKinney, Wes. (2017). Python for data analysis: Data wrangling with Pandas, NumPy, and IPython. O'Reilly Media: Good manual for understanding the capabilities of pandas as well as its strengths written by pandas' creator. Learn how to manipulate, process, clean, and crunch data in Python and learn Pandas for real work
- Sweigart, Al. (2019). Automate the boring stuff with Python: practical programming for total beginners. No Starch Press: Really useful worked examples on how to use Python to automate part of your workflow. For the 'total beginner' is probably an overstatement - well suited even for advanced Pythonistas who want to learn some additional tricks
- Reitz, K., Schlusser, T. (2016). The Hitchhiker's Guide to Python: Best Practices for Development (1st Edition). O'Reilly Media. - Not so much to learn Python, but a set of recommendations and best practices for making your code more "pythonic". A bit dated now.
- VanderPlas, Jake. (2016). Python data science handbook: Essential tools for working with data. O'Reilly Media.
- Matthes, E. (2019). Python crash course: a hands-on, project-based introduction to programming. No Starch Press: Comprehensive guide to Python for beginners.

Software: The course will rely primarily on Python and [GitHub](#).

- To install Python, I recommend downloading and installing the [Anaconda Distribution](#), which will automatically install most of the Python libraries we'll be using
- Sign up for a **GitHub** account if you don't have one already
- Install [GitHub Desktop](#), which is a desktop client that simplifies the use of GitHub and has most, if not all, of the functionality you will need
- Python resources
 - [Python Cheat Sheet](#): A refresher and reference for basic Python commands
 - [Awesome Python](#): A curated list of great Python packages and tools

Cloud computing: AWS Academy, or Google Cloud and Databricks classrooms will be available to train you to leverage data at scale.

Data camp: This class is supported by Datacamp, an intuitive learning platform for data science. Learn R, Python and SQL the way you learn best through a combination of short expert videos and hands-on-the keyboard exercises. Through QMSS you are able to access Datacamp for free.

The syllabus indicates suggested Datacamp modules to complement the lecture material and graded assignments. Of course, feel free to try out other offerings you are

interested in from Datacamp's high quality content or other comparable service providers.

Communications: An Ed Discussions workspace for this course will serve as the primary means of written communication before, during and after class, where students can communicate with each other and with instructors. E-mail will be reserved for official communications only.

III. Course Dynamics

Expectation of Regular Participation and Utilization of Course tools: We will be monitoring student participation and completion of assignments using the corresponding tools throughout the semester. We want to make sure that students are consistently engaged, and if that becomes difficult, that students alert us to their situations.

In preparation for each class: You should have (i) consumed all the curated materials listed in the syllabus for the week; (ii) posted any questions you have on **Ed Discussions** and contributed to answer other students' questions; and (iii) submitted homework assignments for that week using **Canvas**.

During each live class: We will engage in a combination of lecture and live-workshop where we will run through code examples, troubleshoot and answer questions. Starter code for in-class exercises will be made available through **GitHub**. You will need to **bring a laptop to class** to follow along the coding tutorials and examples.

IV. Course requirements

The grade for this course will depend on the fulfillment of three main requirements:

- (i) **Attendance and Class participation (10%):** Students are required to attend and actively participate in class exercises. Participation grades will depend on participation in class and in the online Ed Discussion workspace, where you can ask / answer questions about the topics we are learning.
- (ii) **Data Camp exercises (20%):** The syllabus lists data camp exercises to complete throughout the course. If a Data Camp assignment is not listed as optional below, then it is mandatory and should be completed by Wednesday at midnight each week. These are a completion grade and will generally take you 2-4 hours to complete each week.
- (iii) **Take-home exercises (40%):** There will be 6 or 7 homework assignments, which are expected to be turned in before class each week as instructed on **GitHub**.
- (iv) **Final Exam (30%):** The final examination will require students to generate code to perform common data wrangling operations that are required for the successful completion of most data analysis projects and will use mostly pandas. The exam will be held onsite.

Late Submission Policy: All class assignments are expected to be submitted on the due date. Please note that 10% of the maximum grade will be deducted from the score for every day the assignment is submitted late.

V. Course Outline

Week 1 (1/18)– Course Introduction & Setting up Project Workflow

- I suggest coming to class with python installed already, as we will quickly jump in and review/introduce base python. Depending on your comfort with python, a lot of this may be review to you
- **On your own:**
 - Install Python. I recommend you using the [Anaconda Distribution](#) as described above. The Anaconda Distribution includes most Python libraries we will need this semester. It also comes with Jupyter Notebook and Spyder, two IDEs (Integrated Development Environments) that we will use. Throughout the semester, we will mostly rely on Jupyter Notebooks to share code and complete HW assignments.
 - Sign up for GitHub, and don't forget to respond to this Google [form](#) to be added to the course repo since it's private.
 - Install GitHub Desktop. We will be using this client to interact with GitHub. I've personally found it much easier to use than using command-line Git, and as you will see there is almost no learning curve. If you are confident using the command-line for Git, or have another software you like (i.e., Git Kraken), then feel free to skip this step. Many of you will be newish to GitHub, and GitHub Desktop is so intuitive to use that my hope is that you will start using GitHub on EVERY project you complete after this course. The main drawback of using GitHub Desktop is that most help forums online are for command-line Git users
- **Reading:**
 - [McKinney](#), chapter 2 *Python Language Basics, IPython, and Jupyter Notebooks*
- **Datacamp:**
 - (Optional) To review base Python complete the course [Introduction to Python](#)

PART I: Data Manipulation and Coding

Week 2 (1/25) – Git, GitHub, and Markdown in Jupyter Notebook

- **Reading:**
 - A tutorial on [GitHub Desktop](#)
 - [.gitignore](#)
 - [The Ultimate Markdown Guide \(for Jupyter Notebook\)](#)
 - (Optional) [Hello World](#) - GitHub for the non-programming beginner, for those interested in command-line Git
- **Datacamp:**
 - (Optional) [Introduction to Git](#): Feel free to check out this comprehensive introduction to Git, which emphasizes command-line Git.
- **Homework 1:** Using Jupyter Notebooks and Github
 - Only available on GitHub in the repository **QMSS-GR5072_Spring2023_HW/homework-1-main**
 - Your assignment is saved as a README.md file, which I suggest you read directly on the GitHub website
 - Submit your HW files on GitHub in the following directory:
 - ~/Submissions/LASTNAME-FIRSTNAME
 - You will lose points if you don't submit your HW in the correct place

Week 3 (2/1) – numpy array objects

- **Reading:**
 - [McKinney](#), chapter 4 NumPy Basics: Arrays and Vectorized Computation
- **Datacamp:**
 - [Introduction to numpy](#)

Week 4 (2/8) – Data Wrangling with pandas

- **Reading:**
 - [McKinney](#), chapter 5 *Getting Started with pandas*; chapter 7 *Data Cleaning and Preparation*; chapter 10 *Data Aggregation and Group Operations*
- **Datacamp:**
 - [Data Manipulation with pandas](#)
 - (Optional) [Joining data with pandas](#) - chapter 1 & 2 are the standard operations
- **Homework 2:** Data wrangling using Pandas

Week 5 (2/15) – Functions I: Basic structure and logic

- **Reading:**
 - Sweigart, [Automate the boring stuff with Python, chapter 3. Functions](#)
 - [Some basics of code styling](#)
- **Datacamp:**
 - [Chapter 4 on Loops in Intermediate Python](#)
 - [Python Data Science Toolbox \(Part 1\)](#)
- **Homework 3:** for loops and functions

Week 6 (2/22) – Functions II: Nested and complex operations

- **Reading:**
 - [McKinney](#), chapter 3 *Built-In Data Structures, Functions, and Files*
- **Datacamp:**
 - [Python Data Science Toolbox \(Part 2\)](#)
 - Context managers and decorators in the Datacamp course [Writing functions in Python](#)
 - Utilizing classes from [Software Engineering for Data Scientists in Python](#)
- **Homework 4:** Functions II

Week 7 (3/1) – Functions III: Working with strings and dates

- **Reading:**
 - [McKinney](#), chapter 7, 7.3 *String Manipulation* (p.211-218)
 - Sweigart, chapter 7 [Pattern Matching with Regular Expressions](#)
 - [Python RegEx](#)
- **Datacamp:**
 - [Regular expressions in Python](#)

Week 8 (3/8) – Testing

- **Reading:**
 - Chapter 5 *Testing* from Hillard, Dane. (2020). “[Practices of the Python Pro.](#)” Manning Publications (1st edition). Harvard. (PDF)
 - Reitz, K., Schlusser, T. The Hitchhiker’s Guide to Python, [Testing](#) (p72-81)
 - **Datacamp:**
 - [Unit Testing for Data Science in Python](#)
-

PART II: Python and leveraging other software tools

Week 9 (3/15) – ACADEMIC HOLIDAY

Week 10 (3/22) – Introduction to Relational Databases and Advanced Excel

- **Reading:**
 - TBD
- **Datacamp:**
 - [Data Analysis in Excel](#)

Week 11 (3/29) – Working with SQL

- **Reading:**
 - There are lots of great SQL Tutorials to go further. Here are a few pointers:
 - [Codecademy’s SQL Tutorial](#)
 - [SQLzoo’s SQL Tutorial](#)
 - [W3School.com SQL reference \(also interactive\)](#)
 - [Practice using SQLZOO](#)
- **Datacamp:**
 - [Introduction to SQL for Data Science](#)
 - Joins in SQL are essential, especially given the usually relational nature of the data. Cover joins and relational set theory in [SQL for Joining Data](#)
- **Homework 5:** Practicing SQL Queries

Week 12 (4/5) – Working with APIs

- **Reading:**
 - McKinney, *Interacting with Web APIs* (Section in Chapter 6)
 - [API Integration in Python – Part 1](#)
 - [Getting started with APIs in Python to Gather Data](#)
- **Datacamp:**
 - [Intermediate Importing Data in Python](#)
- **Homework 6:** Strings and APIs

Week 13 (4/12) – Web scraping from HTML

- **Reading:**
 - Sweigart, [chapter 12 Web Scraping](#)
- **Datacamp:**
 - [Web scraping in Python](#)
- **Homework 7:** Web Scraping

Week 14 (4/19) – Cloud computing: Amazon Web Services and Parallelization

- **Reading:**
 - TBD
- **Datacamp:**
 - [Introduction to AWS Boto in Python](#)

Week 15 (4/26) – Review and A brief intro to other miscellaneous topics

- The first half of this lecture will focus on review as you prepare for your final exam
- In the second half, we'll learn about various miscellaneous topics as we have time. These topics will not be included in the final exam:
 - Python and working with:
 - Duplicates
 - Missing data
 - JSON & XML data
 - Multicore processing in Python
 - pyenv and poetry
 - Databricks
 - Airflow and Spark
- **Reading:**
 - [Managing Multiple Python Versions With pyenv](#)
 - chapter 6 Shipping Great Code from Reitz, K., Schlusser, T. The [Hitchhiker's Guide to Python](#)
 - [An Overview of Packaging for Python](#)
- **Datacamp:**
 - (Optional) [Introduction to Airflow in Python](#)
 - (Optional) The course [Introduction to PySpark](#) is beyond what we can cover but provides a good start on Spark using R. More info about the Big Data benefits and use of Spark is discussed in Big Data Fundamentals with PySpark

Week 16 (5/10) – Final Exam

Statement on Academic Integrity

Columbia's intellectual community relies on academic integrity and responsibility as the cornerstone of its work. Graduate students are expected to exhibit the highest level of personal and academic honesty as they engage in scholarly discourse and research. In practical terms, you must be responsible for the full and accurate attribution of the ideas of others in all of your research papers and projects; you must be honest when taking your examinations; you must always submit your own work and not that of another student, scholar, or internet source. Graduate students are responsible for knowing and correctly utilizing referencing and bibliographical guidelines. When in doubt, consult your professor. Citation and plagiarism-prevention resources can be found at the GSAS page on Academic Integrity and Responsible Conduct of Research.

Failure to observe these rules of conduct will have serious academic consequences, up to and including dismissal from the university. If a faculty member suspects a breach of academic honesty, appropriate investigative and disciplinary action will be taken following the Dean's Discipline procedures.

Statement on Disability Accommodations

If you have been certified by Disability Services (DS) to receive accommodations, please either bring your accommodation letter from DS to your professor's office hours to confirm your accommodation needs, or ask your liaison in GSAS to consult with your professor. If you believe that you may have a disability that requires accommodation, please contact **Disability Services** at 212-854-2388 or disability@columbia.edu.

Important: To request and receive an accommodation you must be certified by DS.